

Paying for Your Own Data: A Closed-Loop Benchmark for Scientific Machine-Learning Agents

Hindy Ros¹ and Deepjyoti Deka²

¹ Massachusetts Institute of Technology, Cambridge, MA, USA

² MIT Energy Initiative, Massachusetts Institute of Technology, Cambridge, MA, USA

E-mail: hindyros@mit.edu

Abstract. Benchmarks for coding and machine-learning agents almost always hand the agent its data. We present a benchmark in which the agent must instead *generate its own data* from an expensive physics simulator under a fixed budget of solver calls, and then ship a predictor that is scored independently on a frozen set of cases it never sees. The task is surrogate modelling of the Grad–Shafranov plasma equilibrium: the agent must learn to drive a finite-element solver, design a sampling campaign, train a model mapping five plasma-shape parameters to a 96×64 magnetic-flux field, and assess itself honestly. Every condition emits the same machine-checkable deliverable, so any agent substrate is scored on identical terms.

Across 59 runs on a single language model we find, first, that *giving an agent a proven library makes it worse*: from-scratch agents beat library-assisted ones in all four substrates we tested (median rel- L_2 0.137 against 0.229; stratified rank test $z = -4.17$, permutation $p \approx 5 \times 10^{-5}$, blocked on substrate). A deterministic analysis of the code each agent wrote explains why: given a library, all four substrates settle on its default model and its named acquisition strategies; given nothing, all four converge on a different method family, and two of them beat the library’s own score.

Our central finding concerns verification. **Fifteen of 49 scored runs passed every contract gate while being physically invalid** — a rate that depends strongly on the agent substrate and not at all on how much scaffolding it was given. The failures are not subtle once looked for: one run scored 0.071, better than most of the corpus, while emitting a magnetic field throughout the vacuum region where the field is undefined. Self-reported error diverges from independently measured error by up to 0.93, and the runs that misreport are largely *not* the runs that are invalid, so neither measurement detects the other. Machine-checkable deliverable contracts, the standard instrument in this literature, detect none of it. We therefore separate three questions a benchmark must ask — did the agent deliver, is the artifact physically valid, and did it tell the truth — and measure all three, alongside a fourth: did it build what it said it built.

Keywords: agent benchmarks, large language models, surrogate modelling, active learning, plasma equilibrium, scientific machine learning

1. Introduction

Benchmarks for machine-learning agents have settled into a shape. You give the agent a dataset and a metric, and you see how good a model it can fit. MLE-bench draws its tasks from Kaggle competitions; MLAGentBench, DSbench, ScienceAgentBench and their relatives vary the domain and the scaffolding but keep the shape.‡ In every case the data exists before the agent starts, and it is the same data for every agent. That is exactly what makes those benchmarks clean comparisons.

It is also what makes them unlike most of computational science. There, the expensive step is usually *getting* the data — running the simulation, performing the experiment, doing the solve — and the decision that separates a good researcher from a mediocre one is *where to spend the next unit of budget*. An agent that is excellent at fitting a dataset it was handed may be useless at deciding which experiments are worth running. Nothing in the current crop of benchmarks would tell us either way.

This paper introduces a benchmark that puts that decision back. The agent is given a physics solver, a five-dimensional space of inputs, and a hard budget of solver calls. It must work out how to drive the solver, decide which points in the space to spend its budget on, train a model on what it collected, and hand over a predictor — which we then score, on a set of cases it has never seen, with machinery it does not control (Figure 1).

Why this particular solver. The ground truth comes from TokaMaker [1], a finite-element solver for the equilibrium of a magnetically confined fusion plasma. Section 2 explains the physics from scratch; the properties that make it a good instrument for measuring agents are these. It is *genuinely expensive* — seconds per solve, so a budget bites. It *fails*, and does so as a function of the inputs, so the feasible region has to be discovered rather than assumed. Its output is *high-dimensional* (a field on a 96×64 grid) while its input is five numbers, so dimensionality reduction is a real design decision rather than a formality. It is *unforgiving about physics*: the quantity being predicted is a signed physical field in webers, undefined outside the plasma, so an agent that quietly normalises it, or fills in the region where it does not exist, produces something that looks plausible and is wrong. And the answers are *not memorisable* — these particular equilibria are not on the internet.

What we found. Three things, in increasing order of how much they surprised us.

First, with the language model held fixed across every condition, *which agent framework you use matters more than how much help you give it*. The per-cell medians span a factor of 8.8 across four substrates, against a factor of 1.7 between the two levels of scaffolding.

Second, *the scaffolding helped in the wrong direction*. Agents allowed to import a proven, tested pipeline library did worse than agents made to write everything

‡ Citations are collected in Section 3.




Figure 1. The benchmark loop. Everything inside the dashed line is the agent’s own business: it is handed a frozen task specification and a solver, and it must decide which equilibria to compute, spending a fixed budget across an initial design and up to three adaptive rounds. What it hands over is a small set of files. Everything to the right of the dashed line happens outside the agent’s reach — the files are checked mechanically, the predictor is run on sixty held-out cases the agent never sees, and the result is put through diagnostics that ask whether the field is physically meaningful and whether the agent’s own report of its performance was true. Every condition in this paper, including the two that use no language model at all, passes through that identical right-hand column.

themselves — in all four substrates, with no exceptions. A deterministic analysis of the code they wrote (Section 8) says why, and it is not that the library is bad: given the library, all four substrates adopt its default model and the acquisition strategies its documentation names; given nothing, all four independently find a different method family, and two of them beat the library’s own score on the same test set.

Third, and this is the finding we think matters beyond this benchmark: **31% of the runs that passed every machine-checkable gate were physically invalid.** Not marginal — invalid. One of them scored better than most of the corpus on the error metric while emitting a magnetic field throughout the vacuum region outside the plasma, where the field is not defined at all (Figure 7, second row). The contract that certified it is the standard instrument in this literature: check the files exist, check they parse, check the predictor runs and returns the right shape. It cannot see this class of failure, and neither can the headline error number.

Contributions.

- (i) **A closed-loop scientific-ML benchmark** in which the agent buys its own training data from an expensive solver under a fixed budget, with a machine-checkable deliverable, independent scoring on a frozen set, and a four-rung ladder of scaffolding

from a zero-LLM scripted policy to from-scratch — every rung inside the identical contract (Section 4).

- (ii) **A model-controlled comparison of four agent frameworks.** Every cell runs the same model, so “which framework” is not confounded with “which model” — a confound we had in our own earlier campaigns (Section 6).
- (iii) **A verification layer** that separates delivery, physical validity and self-report honesty, and finds all three come apart (Section 7).
- (iv) **A deterministic, LLM-free account of what each agent actually built** — its chain of methods, its per-round decision logic, and whether its code agrees with its own prose. This is what makes the scaffolding result interpretable rather than merely surprising (Section 8).

2. The problem, for a reader who has never seen a tokamak

This section explains the physics the benchmark is built on. A plasma physicist can skip it; nothing later depends on more than is written here.

2.1. What is being computed, and why anyone wants it fast

A tokamak confines a hot plasma inside a doughnut-shaped magnetic field. For the plasma to sit still rather than drift into the wall, the outward pressure of the gas has to be balanced everywhere by the inward force of the magnetic field. A configuration in which that balance holds is called an *equilibrium*, and finding it is the first thing anyone does when designing a device, planning a discharge, or controlling one in real time.

Because a tokamak is (to a very good approximation) symmetric about its central axis, the three-dimensional problem collapses to two dimensions: a single cross-sectional slice, with coordinates R measured outward from the axis of symmetry and Z measured vertically. On that slice the force balance becomes one elliptic partial differential equation, the Grad–Shafranov equation [2, 3]:

$$\Delta^* \psi = -\mu_0 R^2 p'(\psi) - F(\psi)F'(\psi), \quad \Delta^* \equiv \frac{\partial^2}{\partial R^2} - \frac{1}{R} \frac{\partial}{\partial R} + \frac{\partial^2}{\partial Z^2}. \quad (1)$$

The unknown $\psi(R, Z)$ is the *poloidal flux function*, measured in webers. Two facts about it are all that the rest of this paper needs.

The contours of ψ are the surfaces the plasma lives on. Magnetic field lines lie within surfaces of constant ψ , so a plot of ψ is a picture of the magnetic structure: a set of nested closed loops around a central peak. Figure 2(b) is a real solution. The outermost closed contour is the plasma boundary, conventionally the *last closed flux surface* or LCFS.

Outside that boundary, ψ is not defined here. We solve the *fixed-boundary* problem: the boundary shape is prescribed, and the equation is solved on the region inside it. There simply is no value at grid points outside, and the solver returns NaN there. About 80% of the grid is NaN in a typical case. This sounds like a bookkeeping

detail. It is the single most consequential convention in the benchmark, and Section 7 is largely about what agents do with it.

Why a surrogate model is worth having. A finite-element Grad–Shafranov solve takes seconds. That is fine once, and ruinous inside a loop. Device design searches over shapes; scenario planning sweeps over operating points; real-time shape control needs an equilibrium every millisecond, and the iterative solve is precisely the latency bottleneck that prevents higher-bandwidth control. Learning a fast approximation of the solver — a *surrogate* — is therefore an active line of work in its own right, with recent efforts targeting millisecond-scale neural-operator equilibria [4] and control-oriented surrogates validated on an operating device [5]. The task we set the agents is a scaled-down version of a problem people are actually paid to solve.

2.2. The five numbers that define a case

The boundary is a D-shape, the standard cross-section of a modern tokamak, described by four geometric parameters plus the current flowing in the plasma (Figure 2(a)):

Symbol	Name	Range	What it does
r_0	major radius	0.35–0.55 m	how far the doughnut’s centre sits from the axis
a	minor radius	0.10–0.20 m	how fat the doughnut is
κ	elongation	1.0–1.6	vertical stretch; 1 is a circle
δ	triangularity	0.0–0.4	pulls the top and bottom inward into a D
I_p	plasma current	80–200 kA	sets the overall magnitude of ψ

Figure 2(c) shows what the two shape parameters do. The task, then, is to learn the map

$$(r_0, a, \kappa, \delta, I_p) \longmapsto \psi(R, Z) \in \mathbb{R}^{96 \times 64}, \quad (2)$$

from a budget of solver calls the agent itself chooses.

2.3. How good is a prediction? Reading the error metric

Predictions are scored by relative L_2 error, computed per case over the grid points where the truth is defined:

$$\text{rel-}L_{2i} = \frac{\|\hat{\psi}_i - \psi_i\|_2}{\|\psi_i\|_2}, \quad \text{over } \{(R, Z) : \psi_i(R, Z) \text{ finite}\}. \quad (3)$$

Zero is perfect. The number that gives it meaning is what you get for free: a predictor that ignores its input entirely and always returns the average flux map of the training set scores 0.4933 on our test set. We call that the *mean-map baseline* and quote accuracy against it, $100 \times (1 - \text{rel-}L_2/0.4933)$, so that 0% means “no better than not trying” and 100% means exact.

Because an error metric is easier to trust once you have seen it, Figure 3 shows the same equilibrium predicted at a range of rel- L_2 values by actual runs from our campaign.




Figure 2. (a) The plasma boundary and the four shape parameters that set it; the fifth, the plasma current I_p , scales the field rather than the shape. The dashed line at $R = 0$ is the machine’s axis of symmetry, about which the whole picture is revolved to make the torus. (b) A real solution from the frozen test set. Colour is the poloidal flux ψ ; the black contours are flux surfaces, and the heavy line is the plasma boundary. Outside it, ψ is undefined and the solver returns **NaN** — roughly 80% of the grid. (c) What elongation and triangularity do to the boundary: κ stretches it vertically, δ pulls the top and bottom inward.




Figure 3. What the error metric looks like. The same test equilibrium, predicted by six actual runs spanning the range we observed, with rel- L_2 for that case above each panel and the condition that produced it below. “Mean map” is the input-independent baseline. All panels share a colour scale.

Below about 0.1 the prediction is visually indistinguishable from the truth. By 0.2 the field is noticeably blurred. At the mean-map baseline it is a smear with no case-specific structure at all, and beyond that the predictor is contributing nothing.

3. Related work

Benchmarks that hand the agent its data. The dominant shape gives an agent a repository or a dataset and scores the artifact it produces. SWE-bench [6] scores patches against held-out tests; AgentBench [7] covers interactive environments more broadly. Closest to our setting are the ML-engineering suites: MLE-bench [8] runs agents on Kaggle competitions and grades against the private leaderboard; MLAGentBench [9] scores iterative experimentation on fixed task datasets; DSbench [10] draws on data-science competitions.

A second line targets science specifically. ScienceAgentBench [11] assembles data-driven discovery tasks from published papers; DiscoveryBench [12] formalises hypothesis search over provided datasets; BixBench [13] does the same for computational biology; CORE-Bench [14] scores computational reproducibility; PaperBench [15] scores replication of ML papers; RE-Bench [16] pits agents against human experts on open-ended research engineering.

These differ enormously in domain and open-endedness, and in one respect they agree: the training data exists in advance and is identical for every agent. That is what makes them clean comparisons — and it is also what removes acquisition from the set of decisions being measured.

The literature the task asks an agent to rediscover. Deciding where to measure next is of course a mature field. Active learning [17] and Bayesian experimental design [18] supply the strategies a competent agent might reach for — uncertainty sampling, residual-targeted acquisition, variance reduction under a probabilistic model — and space-filling designs such as Latin hypercubes [19] supply the baseline those strategies have to beat. We contribute nothing to that literature. What we measure is whether an agent reaches for any of it unprompted, implements what it says it implemented, and spends a hard budget sensibly (Section 8).

What we add, stated narrowly. The gap is not that nobody acquires data. It is that *agentic coding benchmarks do not make the acquisition policy part of what is scored*. To our knowledge, as of mid-2026, no existing agent benchmark requires the agent to (i) drive an expensive, failure-prone numerical solver for which it has not been handed a working harness, (ii) allocate a hard budget across adaptive rounds, and (iii) be scored on a frozen set it never sees, with a zero-LLM scripted policy running inside the identical contract as the reference point.

Two recent works come closest. AstaBench [20] assembles a large scientific-agent suite and is motivated by the same complaint about confounders — model cost and tool access varying across compared systems — that motivates our model control; it supplies the tools and the data rather than scoring acquisition. And a position paper argues that the execution harness, not the model it wraps, is often the binding constraint on agent performance [21]. Our campaign is a direct empirical test of that claim: one model, four harnesses, and the spread measured (Section 6.3).

The agent frameworks we evaluate. These are existing systems, not contributions of this paper: URSA [22], a planner/executor pair built on LangGraph, part of the LangChain project [23]; DSPy [24], whose typed modules we also use for the L1 decision layer and whose GEPA optimiser [25] we use for prompt optimisation; and two commercial command-line coding agents. We evaluated them as we found them; none of their authors was involved in or consulted about this evaluation.

4. The benchmark

4.1. Vocabulary

Five terms recur, and it is worth fixing them before they are used.

Term	Meaning in this paper
<i>harness</i> (or <i>substrate</i>)	The agent framework: the software that wraps a language model and turns it into something that can read files, run commands and iterate. We compare four. It is not the model.
<i>access level</i>	How much pre-written help the agent is given, from L0 (a scripted policy, no model at all) to L3 (write everything yourself). Section 4.4.
<i>the contract</i>	The fixed set of machine-checkable conditions every condition’s output must satisfy: the files exist, the report parses, the predictor runs and returns the right shape on the right grid. Section 4.3.
<i>the frozen set</i>	Sixty solved equilibria, fixed once and never regenerated, on which every predictor in this paper is scored. No agent ever sees them.
<i>honesty gap</i>	An agent’s own reported error minus its independently measured error. Negative means it claimed to be better than it was. Section 7.4.

4.2. What the agent is asked to do

Each agent receives a frozen, versioned problem statement — the same text for every replicate, hashed into every run record, and reproduced verbatim in Appendix Appendix D. It specifies five things.

The forward problem. Fixed-boundary equilibria of D-shaped plasmas in the five parameters of Section 2, over the stated ranges, at the solver’s default profile shapes.

A campaign structure with a hard budget. At most 150 solves for an initial space-filling design, then up to three adaptive rounds of exactly 100 new solves each. Failed solves count against the budget. Crucially, each round’s points must be chosen from the current model or the data collected so far: a single large one-shot design is explicitly ruled out, because the object of study is the choice of where to sample.

A stopping criterion. Keep a validation set separate from the test set, and stop early once validation error falls to $0.30\times$ the mean-predictor baseline or better.

An honest evaluation requirement. Hold out at least 20 solves at random parameters with a recorded seed, and never use them for training, model selection or acquisition.

A self-assessment requirement. Report whether the adaptive sampling actually helped, “or an honest statement that it did not”.

Four process gates are mandatory and identical for every condition: a *meshing milestone* (reproduce the solver’s own worked example before generalising), a *pilot gate* (at least 20 solves at 50% success before the rounds begin), a *storage-validation gate* (a solve counts as successful only once its stored file has been read back and confirmed to contain finite numbers), and a *deliverable self-test* (run every documented command in a fresh process before declaring done). The last two were added after an earlier campaign

Table 1. The deliverable contract. The scoring subprocess is stripped of provider credentials, so `predict.py` must work offline, and it runs in its own process group so a timeout reaps the solver’s children too. The final gate applies only at L3, so L3 runs face nine checks and L2 eight.

Gate	What it checks
<code>artifact:report.json</code>	present and non-empty
<code>artifact:predict.py</code>	present and non-empty
<code>artifact:README.md</code>	present and non-empty
<code>report_parsing</code>	<code>report.json</code> is valid JSON
<code>report_keys</code>	the required keys are <i>present</i>
<code>predict_runs</code>	<code>python predict.py --input p.json --output o.npz</code> exits 0
<code>predict_shape</code>	the returned <code>psi</code> is $(N, 96, 64)$
<code>predict_grid</code>	the returned <code>R</code> , <code>Z</code> match the frozen axes to 10^{-9}
<code>no_autotokamak_import</code>	L3 only: no workspace file imports the platform library

in which an agent stored all-NaN data, counted it as success, and shipped a predictor it had never executed.

4.3. The deliverable contract

Every condition emits the same three files — `report.json`, `predict.py`, `README.md` — and they are checked mechanically (Table 1). `contract.passed` is the flat AND of every gate: no weighting, no partial credit.

One design choice in that table does a great deal of work later. `report_keys` **checks that the agent’s self-reported metrics are present, and never that they are correct**. The agent’s own numbers are recorded verbatim and never cross-checked by the contract. That is what makes the honesty measurement of Section 7.4 possible at all: we can compare what the agent said about itself against what we measured, precisely because nothing in the pipeline ever forced the two to agree.

Predictions are scored on the frozen set of 60 parameter vectors described in Section 2, using Equation (3). One detail matters: **a NaN prediction at a point where the truth is defined is scored as a full-magnitude miss**, never skipped. Skipping would let an all-NaN prediction score as perfect.

4.4. The scaffolding ladder

The benchmark’s second axis is how much help the agent gets. All four rungs emit the identical contract and are scored on the identical frozen set, which is what makes them comparable at all.

Level	Name	What the agent does
L0	scripted	Nothing. A seeded heuristic policy drives a pre-written pipeline; no language model is involved at any point.
L1	structured	A model makes only typed decisions — which model family, how many components, what to do next — inside fixed pipeline code.
L2	library-assisted	The agent writes the glue code and <i>may</i> import the platform library, which already implements the sweep, the acquisition strategies and the model zoo.
L3	from-scratch	The agent writes everything. Importing the platform library is forbidden, and the prohibition is audited by static analysis of every file it wrote.

L0 is the reference point that makes the rest interpretable. It is deterministic, it costs nothing, and it is scored by the same scorer on the same cases. Without a zero-model condition inside the same contract there is no way to say how much of a given score is the task being easy and how much is the agent being good.

4.5. Is the task hard for the right reasons?

Before asking whether agents can do this, it is worth establishing that the task rewards the thing we claim it rewards. Three facts, established by a direct sweep of 9993 converged solves independent of any agent, are shown in Figure 4; the underlying tables are in Appendix Appendix C.

First, error falls with training-set size and saturates — and *the best model family changes along the way*, from polynomial ridge up to $N = 200$ to kernel ridge from $N = 400$ on. Model selection is therefore a real decision that depends on the data the agent chose to collect, not a fixed answer that can be looked up. The agents’ 450-solve budget sits in the part of that curve where it still matters.

Second, at a matched budget, *where* you sample beats *how much*. Space-filling designs beat uniform random by 21% at $N = 100$, narrowing to 6% by $N = 1600$; deliberately clustered sampling is catastrophic at every budget. Coverage is the first-order concern and its value decays as the budget grows.

Third, and this is what makes the sampling decision consequential rather than cosmetic: *outside its training region the surrogate collapses*. Train on half the parameter space and test on the other half and the error ratio is 0.94 — no better than predicting the mean field — against 0.38–0.40 within distribution. An agent that samples badly does not lose a few percent. It produces a model that fails completely off-distribution and, because it evaluates itself on its own test set drawn from its own sampling distribution, it may never find out.




Figure 4. The task is hard for the right reasons. **(a)** Error against training-set size, with the best model family at each size marked; the shaded band is the agents’ campaign budget. **(b)** At matched budget, sampling strategy dominates. **(c)** Trained on one half of the parameter space and tested on the other, the surrogate is no better than the mean map. All values are relative L_2 divided by the mean-map baseline, from a pool of 9993 converged solves.

4.6. A worked example

To make the loop concrete, here is one run end to end: L3-cursor replicate 2, chosen because it is unremarkable — it sits in the middle of its cell and nothing went wrong.

It received the task text of Section 4.2 and a workspace containing nothing but a read-only copy of the solver’s source tree. It began, as instructed, by locating the solver’s own fixed-boundary example and reproducing it as a standalone script (`meshing_milestone.py` in its workspace), confirming it could produce a finite flux field before writing anything of its own. It then built a mesh-and-solve wrapper, ran a 20-solve random pilot to check the failure rate against the mandated 50% threshold, and extended to 150 attempted solves with a Latin hypercube over the five parameters. It set aside a further 30 solves for validation and 25 for its own held-out test, and recorded a separate seed for each of the four.

For the model it chose something no prompt suggested: a PyTorch multilayer perceptron mapping the five parameters directly to the 96×64 field, trained under a loss masked to the points inside the plasma. It left dropout active at inference so that repeated forward passes give a spread, and used that spread as its acquisition signal — sampling where the dropout ensemble disagreed most. One adaptive round of 100 solves was enough: validation error crossed the task’s $0.30 \times$ -baseline threshold and it stopped early, having used 305 of its 450 permitted solves.

It wrote the three required files, re-ran its own `predict.py` in a fresh process as instructed, and declared itself done after 8.9 minutes and about \$1.64 of model spend. Our independent scoring found rel- L_2 0.0622 — 87.4% accuracy against the mean-map baseline, the fourth-best run in the campaign. All nine gates passed, all four physical-

validity diagnostics passed, and its NaN mask agreed with ground truth on 99.98% of grid points. Its own report claimed 0.0611 against our 0.0622: a gap of -0.001 , which is to say it flattered itself by a tenth of a percent. It also volunteered, unprompted, that its adaptive round had helped — 0.085 before, 0.061 after.

That is what the benchmark looks like when everything works. Much of the rest of this paper is about the runs where it did not, and about how little the contract noticed.

5. Experimental design

5.1. Conditions

The campaign is the cross of two access levels and four agent frameworks: $\{\text{L2}, \text{L3}\} \times \{\text{ursa}, \text{dspace}, \text{pi}, \text{cursor}\}$, eight cells, replicated — plus a budget-matched L0 reference arm. We attempted 59 runs and scored 49.

Replicate counts are deliberately unequal: two substrates at $n = 10$ and two at $n = 5$. The campaign was budgeted in dollars rather than in runs, and one substrate costs an order of magnitude more per run than another (Section 6.5). Failures reduce some cells further. Every denominator in this paper counts *attempts*, not survivors, which is the conservative choice and is stated wherever a rate appears.

5.2. Holding the model fixed

Every cell runs the same language model, `gpt-5.2`. This is the point of the design. In an earlier campaign of our own, one harness ran a different model family from the others, and the resulting comparison was uninterpretable: any difference could have been the framework or the model, and nothing in the data could separate them.

Pinning it is fiddlier than it sounds, because the four adapters disagree about how a model is named — two take `openai:gpt-5.2`, one takes `openai/gpt-5.2`, and the command-line agent takes a bare `gpt-5.2`. We pin the name per adapter inside the frozen task specification rather than passing a single command-line override, so that the model string is recorded three independent times per run and the control is auditable after the fact. A fifth adapter we maintain is excluded from the campaign entirely for the same reason: it cannot run the pinned model, and including it would have reintroduced exactly the confound the design exists to remove.

Framework versions were `ursa` 0.15.1, `dspace` 3.2.1, `pi` 0.73.1 and `cursor` 2026.09.18, against Open FUSION Toolkit 26.6. Those are recorded from the machine, not from the run records — which stamp the task hash, the prompt version, the model and the replicate index, but not the substrate’s own version. That is a gap in our harness and we would close it in a future campaign.

5.3. Equal compute budgets, and why they are hard to enforce

Giving every substrate the same budget turned out to be the hardest engineering requirement in the benchmark, and the failure modes generalise well past this paper.

Two adapters accepted a `timeout_seconds` argument and silently ignored it. Fixing that exposed two deeper problems. First, **a timeout raised as an ordinary exception is swallowed**: agent frameworks wrap their step loops in broad `except Exception` handlers and treat anything caught as a recoverable step failure, so a timeout delivered by signal vanishes into the agent’s own retry logic — and because `signal.alarm` is one-shot, a single swallowed alarm disables the cap permanently. One harness absorbed its 90-minute cap this way and ran for over ten hours. Second, **killing the agent process does not stop the work**: Python’s `subprocess.run`, called with a timeout, kills the direct child and then waits on its pipes again — but the solver’s grandchildren inherit those pipes and hold them open, so the timeout path itself blocks. One cell sat in that deadlock for 9 hours 19 minutes.

The fixes are to raise a `BaseException` subclass, for the same reason `KeyboardInterrupt` is one, and to run each agent in its own process group and signal the group. The general lesson is worth stating plainly: **an agent harness is third-party code and cannot be trusted to honour a deadline from inside**. We enforce the budget at three independent levels — in-adapter, by process-group kill, and by a driver-side watchdog that nothing running inside the cell can intercept.

We also equalised the outer feedback-round setting to 1. Only two of the four adapters implement outer review-and-fix rounds at all, so any higher value hands those two a second pass the others never get — an arbitrary two-tier design that was the mechanical cause of their longer runtimes and larger spend in earlier campaigns. One is the only value all four honour identically.

5.4. What is held fixed

The prompt text (version 3, hashed per run), the evaluation grid, the 60-parameter test set, the gate set and the scoring rule. The L2 and L3 problem statements are byte-identical except for ten passages, every one of which concerns access level; they are reproduced with the differences marked in Appendix Appendix D. Results are pooled only within one prompt version and one scoring epoch, and the aggregation tool refuses to mix them.

5.5. Analysis, and one change we made after seeing the data

The unit of analysis is the cell. Pass rates carry Wilson score intervals [26], which — unlike the normal approximation — do not collapse to zero width at $0/n$ or n/n . Relative- L_2 is summarised by medians with percentile bootstrap intervals [27] rather than means, because a single catastrophic replicate makes a mean uninterpretable and we have several.

The L2/L3 contrast is paired within harness, since the harness is a blocking factor rather than a treatment. We pre-specified an exact two-sided sign test over the four harness pairs.

That test cannot return a p below 0.125 at four pairs, whatever the data show, and the campaign produced a unanimous 4/4 direction. We therefore **changed the primary test after seeing the data**, to a stratified rank statistic in the van Elteren family [28], which uses all 49 scored runs while still permuting labels only within harness. This is a post-hoc change and we treat it as one: the analysis tool prints all three tests on every run, all three are reported in Section 6.2, and the conclusion carries the caveat. Per-cell orderings are exploratory throughout; the pooled level contrast is the claim that carries weight.

6. Results

The campaign ran 59 cells to completion or failure and scored 49 of them, at a measured cost of \$197.78.§ Ten runs produced no score at all; Section 7.5 accounts for each of them, and none is a failure of the model’s ability to do the task.

6.1. The per-cell table

Table 2 is the campaign in one view: eight agent cells, the zero-model reference arm, and for each of them how often the deliverable passed, how good it was, how much of it was physically valid, and what it cost. The sections that follow take one column of it at a time.

6.2. Does library access help? It does not

Pooled across substrates, L2 has median $\text{rel-}L_2$ 0.2287 (bootstrap [0.211, 0.396], 24 scored runs) against L3’s 0.1368 ([0.077, 0.207], 25 runs). Contract pass rates are 24/33 (Wilson [0.56, 0.85]) and 25/26 ([0.81, 0.99]). Accuracy per hundred solver calls — a crude efficiency measure, but one that accounts for agents that bought their way to a better score — is 10.31 against 19.55.

The paired differences are -0.149 , -0.216 , -0.083 and -0.376 for `cursor`, `dspy`, `pi` and `ursa` respectively: four substrates, four times in the same direction (Figure 5).

Our primary test treats the harness as a blocking factor and permutes level labels only within it — a stratified rank statistic in the van Elteren family [28]. It gives $z = -4.168$ with permutation $p = 1/20\,001 \approx 5 \times 10^{-5}$. That p is the resolution floor of the test rather than a point estimate: zero of 20 000 within-harness relabelings reached the observed statistic. Substrate differences cannot inflate it, because no permutation ever moves a run between harnesses.

Two other tests, reported because they differ in power rather than in direction. The exact sign test on the four paired medians gives 0/4 and $p = 0.125$ — the smallest two-sided p attainable at four pairs, whatever the data show. The originally pre-specified

§ From the directory-keyed cost report. Our aggregation tool reports \$199.76 because it joins costs on a run identifier that two substrates reuse across sibling replicate directories; we report the former and note the defect rather than silently choosing between them.

Table 2. Per-cell results. n counts *attempts*, so a replicate killed before it could write a deliverable counts against the pass rate; $sc.$ counts replicates that produced a frozen-set score. Relative L_2 is the median over scored replicates with a percentile bootstrap interval (10000 resamples, seed 20260917); the mean-map baseline is 0.4933. “Valid” passes every gate *and* the physical diagnostics of Section 7.1; in every cell the complement is a run that passed every gate while being physically invalid. *cursor* and *ursa* costs are derived from token counts at a committed price table rather than vendor billing.

Level	Harness	n	$sc.$	pass k/n [Wilson]	med.	bootstrap CI	valid	\$
L0	scripted*	3	3	n/a	0.1863	0.1774–0.1865	n/a	0.00
L2	<i>cursor</i>	5	5	5/5 [0.57, 1.00]	0.2186	[0.175, 0.234]	5/5	5.12
L2	<i>dspy</i>	15	10	10/15 [0.42, 0.85]	0.4146	[0.263, 0.503]	5/10	21.37
L2	<i>pi</i>	5	5	5/5 [0.57, 1.00]	0.1878	[0.167, 0.211]	5/5	3.94
L2	<i>ursa</i>	8	4	4/8 [0.22, 0.79]	0.6101	[0.218, 2.494]	2/4	50.71
L3	<i>cursor</i>	5	5	5/5 [0.57, 1.00]	0.0692	[0.062, 0.212]	5/5	9.23
L3	<i>dspy</i>	10	10	10/10 [0.72, 1.00]	0.1985	[0.119, 0.229]	4/10	23.46
L3	<i>pi</i>	5	5	5/5 [0.57, 1.00]	0.1052	[0.062, 0.130]	5/5	6.16
L3	<i>ursa</i>	6	5	5/6 [0.44, 0.97]	0.2342	[0.019, 1.084]	3/5	77.79

* the L0 arm is the pipeline run directly, not an agent in a workspace: it emits no deliverable contract, so the gate and diagnostic columns do not apply. Its interval is the range over three seeds, not a bootstrap. It is budget-matched to 450 successful solves, makes zero model calls, and is scored by the identical scorer on the identical cases.




Figure 5. Every scored run, by cell. Circles are individual replicates, the heavy bar is the cell median with its bootstrap interval, and the arrow is the within-harness L2→L3 contrast, annotated with the difference in medians. From-scratch is better in all four substrates. The two substrates with the widest replicate spread are also the two that produce runs worse than predicting the mean field.

median-difference permutation test gives $\delta = -0.206$, $p = 0.179$; it compresses each cell to a single number and then has four numbers to work with. As set out in Section 5.5, the switch of primary test was made after seeing that a unanimous result could not clear the sign test’s floor, and is post-hoc.

Two readings survive this. Library access might genuinely handicap an agent, or

the L2 prompt’s pointers into the library might anchor agents onto its defaults. Section 8 shows the anchoring is real and measurable, and Section 6.4 bounds how much of the gap the library’s own quality could explain.

6.3. Does the harness matter, with the model held fixed?

Yes — and more than the access level does. With one model across all eight cells, the per-cell medians span 0.0692 to 0.6101, a factor of 8.8, against the $1.7\times$ pooled level effect. Ordering the substrates by L3 median gives **cursor**, **pi**, **dspy**, **ursa**; the same order holds at L2 but for **pi** and **cursor** swapping places.

The spread is not only in the median, and this is the part we did not expect. **cursor** and **pi** produced twenty scored runs between them with *zero* physically invalid artifacts and a worst case of 0.234. **dspy** and **ursa** account for all fifteen invalid artifacts and every run that scored worse than predicting the mean field. Same model, same task text, same budget; the difference is the software wrapped around the model.

This is the comparison the literature usually confounds, and it is the empirical form of a claim recently argued on other grounds — that the harness, not the model it wraps, is often the binding constraint [21].

6.4. Against a baseline that uses no model at all

The pre-written pipeline, run at L0 with seeded heuristic decisions and no language model anywhere in the loop, scores mean rel- L_2 0.1774, 0.1865 and 0.1863 on three seeds at a budget matched to the agents’ 450 successful solves. That is about 63% accuracy, in 83–101 seconds, for \$0.00 of model spend. Given five times the data it reaches 0.1666, so its budget is not what limits it.

Against that line, **two of the eight agent cells beat the zero-model baseline** — L3-cursor and L3-pi — one roughly matches it, and five are worse. Both that beat it are from-scratch.

This constrains how Section 6.2 may be read. The library is not simply bad code dragging L2 down: it outperforms five of the eight agent cells on the same frozen set. Whatever is happening at L2 is better explained by what library access does to the agent’s search over methods than by the quality of the library itself.

6.5. Cost, and whether the budget was respected

Cost is dominated by a single substrate. **ursa** spent \$128.51 across 14 runs — 65% of the entire campaign — against \$44.82 for **dspy** (25 runs), \$14.36 for **cursor** and \$10.10 for **pi** (10 runs each). It buys nothing legible (Figure 6).

Costs are measured directly where a substrate reports dollars and otherwise derived from token counts at a committed, dated price table. The **cursor** figure is an estimate at provider rates rather than that vendor’s actual billing, because that agent consumes its own product’s quota; it is labelled derived everywhere it appears.




Figure 6. What each cell’s accuracy cost, at one model. A $13\times$ spread in dollars buys no consistent advantage: the two cheapest substrates occupy the two best positions, and the most expensive cell in the campaign sits below L3-pi at one-thirteenth the price.

On the 450-solve campaign bound, the median cell is compliant once the uncapped validation and test solves are allowed for (medians of 239–689 attempted solves). The exceptions are concentrated in L2-dspy, whose worst runs attempted 1914 and 1650 solves — overruns of +1464 and +1200 against a bound stated plainly in the prompt they were given. We report the raw distribution rather than a pass/fail, since validation and test solves are legitimately additional, but an overrun of four times the campaign budget is not an accounting subtlety.

Wall-clock time is contaminated by running several cells concurrently on one CPU-bound machine, and is not a reportable efficiency metric here. For completeness, the campaign consumed 19.1 hours summed over 59 runs.

7. Verification: did it deliver, is it valid, did it tell the truth?

Three questions get collapsed into one in most of this literature. The contract of Section 4.3 answers the first. It cannot answer the other two, and this section is about what that costs.

7.1. Passing every gate is not the same as being right

`contract.passed` certifies *delivery*. It says the files are there, the report parses, the predictor runs and returns an array of the right shape on the right grid. It says nothing about whether that array is a physically meaningful magnetic field.

So we record diagnostics alongside the contract — never as gates. Adding a gate would change `contract.passed` and silently make every previously archived run

incomparable, which is precisely the kind of quiet breakage the benchmark exists to detect. Validity is judged on *magnitude*, as the conjunction of four checks:

- **Sanity:** is $\text{rel-}L_2$ below 0.8? That is, is this a surrogate at all, given that the mean map scores 0.49?
- **Full-grid error:** the same metric recomputed with the plasma’s *exterior* included in the numerator while the denominator stays the interior norm. A spurious field written where flux is undefined is charged for rather than masked away.
- **Exterior inflation:** the ratio of the two. About 1.0 is clean; about 1.2 is a NaN-convention slip; 17 is an invented plasma.
- **Prediction spread:** does the prediction vary across test inputs at least a tenth as much as the truth does? A constant predictor is finite, correctly shaped, and ignores its input.

NaN-mask agreement — whether the prediction’s undefined region matches the truth’s — is recorded but deliberately *demoted* out of the validity rule. Counting pixels cannot distinguish a harmless convention slip from an invented plasma: both can land at 0.195 agreement while differing by an order of magnitude in how wrong the field actually is. Magnitude can tell them apart, so magnitude decides and the mask is kept as evidence.

A run that passes every gate and fails the magnitude rule is recorded as `passed_gates_but_invalid`.

7.2. How often that happens

Fifteen of the 49 scored runs — 31% — passed every contract gate while being physically invalid. The rate is strongly substrate-specific and almost completely level-independent:

	cursor	pi	dspy	ursa	L2	L3
passed every gate, physically invalid	0/10	0/10	11/20	4/9	7/24	8/25

Two substrates produce none of it and two produce all of it, while the axis the benchmark was built to vary — how much scaffolding the agent was given — makes no difference whatsoever. An evaluation reporting only `contract.passed` would have scored this campaign as 49/49 delivered and seen none of this.

7.3. What it looks like

Figure 7 is the whole argument in one image. Four predictors are shown against the same ground-truth equilibrium.

The second row is the one worth dwelling on. That run, an L3-dspy replicate, scored $\text{rel-}L_2$ 0.071 on the interior — 85% accuracy, better than most of the corpus, seventh-best of 49. Its NaN-mask agreement is 0.195: it emitted *no* NaN at all, against a ground truth that is 80.5% NaN. Its full-grid error is nearly double its interior error. It passed all nine gates.




Figure 7. Four predictors against the same test case. The top row is the truth and, beside it, the region where ψ is defined at all. Each subsequent row is one run's prediction and its signed difference from the truth. $\text{rel-}L_2$ is for this case. Residual panels are scaled per row — a scale shared with the bottom row would render the others uniformly blank — with the peak deviation printed on each. **The second row is the point of the figure:** it scores better than the two rows beneath it on the interior metric, and it has painted a magnetic field across the entire vacuum region, where there is no field to predict.

Nothing in the contract could catch this, and neither could the headline number, because the headline number is computed only where the truth is defined. The failure is invisible to both and obvious to the eye.

The same family of failure recurs. An earlier campaign produced a predictor that emitted zeros instead of NaN outside the boundary and passed 9/9 gates at rel- L_2 0.96, twice as bad as predicting the mean field. In this campaign, the third row of Figure 7 has a near-perfect mask and a prediction spread of 0.05 — it has learned the shape of the answer and almost nothing else — and the fourth has the right mask with a field five times too large. All of it is caught by three numbers computed from predictions the benchmark already had in hand.

7.4. The honesty gap

Because the contract checks that self-reported metrics are present but never that they are true, the difference between what an agent says about itself and what we measure is directly observable:

$$\text{gap} = \overline{\text{rel-}L_2}^{\text{self-reported}} - \overline{\text{rel-}L_2}^{\text{measured}}, \quad (4)$$

where negative means the agent claimed to be better than it was.

Across the 48 scored runs that reported a comparable number, the median gap is -0.0012 . **Most agents report their own error accurately**, and that is worth saying clearly before the exceptions.

The tail is another matter. Ten runs exceed $|\text{gap}| = 0.1$. The worst is an L2-ursa run reporting 0.0515 against a measured 0.9857 — a claim of roughly nineteenfold better accuracy than it achieved. In the other direction, one run reported 4.2557 against a measured 2.4941, overstating its own error. The plausibility filter we apply before comparing self-reports, which rejects anything outside $[0, 5]$, admits both of those extremes and so separates nothing in this campaign.

The two pathologies are largely different runs. Of the fifteen physically invalid artifacts, only four have $|\text{gap}| > 0.1$; of the ten large-gap runs, only four are invalid. An agent can be wrong and honest about it, or approximately right and still misreport. *Neither measurement substitutes for the other.* Both cluster in the same two substrates, but within those substrates they mark different replicates. Figure 8 plots the two axes against each other; every point in it passed every gate.

7.5. What went wrong in the runs that produced nothing

Ten of 59 runs produced no score, and it is worth being clear that none of them is a failure of the model’s capability at the task:

- **Five L2-dspy** runs were terminated by the driver watchdog before they could write a deliverable. They are reconciled as explicit non-passing stubs rather than quietly dropped, which is why that cell’s denominator is 15.




Figure 8. The verification plane. Every run shown passed every contract gate. Crosses are the physically invalid ones; the dashed line is the mean-map baseline. Invalid artifacts are scattered across the whole range of self-reported honesty, which is what makes the two measurements independent rather than redundant.

- **Three** L2-ursa runs died after a single turn on one transient provider connection error, at about \$0.05 each, because that adapter has no retry that survives it.
- **One** L3-ursa run hit a timeout.
- **One** L2-ursa run completed with 150 solves attempted, one succeeded, and nothing usable.

Among runs that *did* score, the substantive failures are the fifteen invalid artifacts, and they split by mechanism in a way that a single scalar error could never separate. *Masking* failures have a wrong NaN convention and exterior inflation of 2–4× — the **dspy** pattern, whose modal masking rule is a validity mask derived from training data rather than the plasma boundary itself. *Scale* failures have a near-perfect mask and the wrong magnitude — the **ursa** pattern. Section 8 traces both back to identifiable choices in the code.

7.6. A blind code review, and where it disagrees with the parser

As a qualitative cross-check we ran a blind, outcome-blind rubric review over three replicates per cell: seven dimensions scored 1–5, the median of three passes, 23 judgeable runs, \$2.53 in total. The reviewer sees an anonymised code bundle with brand tokens redacted, never learns which substrate wrote it, and never sees the gates or the frozen score.

It ranks the L3 substrates **pi** 3.71, **cursor** 3.57, **dspy** 3.14, **ursa** 2.71 — the frozen-score ordering, up to the **pi**/**cursor** swap — and rates L3 above L2 overall (3.07 against 2.86).

Where it *disagrees* is the interesting part. It judged the adaptive sampling “genuine”

in 21 of 23 runs, while the deterministic code analysis of Section 8 finds the acquisition criterion actually reacting to what was measured in 2 of 37. A careful reader of the code accepts nominal adaptivity that a parser does not — which is an argument for the parser. The reviewer is also the same model that produced every run it grades, so its scores are used only to corroborate the deterministic analysis and never as a measurement in their own right (Section 9).

7.7. Was the adaptive sampling worth asking for?

The benchmark requires agents to sample adaptively, so it ought to say whether that requirement is justified.

We ran the pipeline’s own meta-loop twice over with the action *forced*, so that the two arms differ in nothing but how new solver calls are chosen: treatment forces residual-UCB acquisition, control forces a blind space-filling append. Four paired replicates, identical per-iteration solve budgets, an envelope evaluation set, and no model calls in either arm. The pre-registered headline is **worst-cell** accuracy rather than aggregate accuracy, on the grounds that an adaptive method ought to pay off precisely where the surrogate is weakest, and an aggregate can hide that.

It is a clean negative. Aggregate accuracy improves by +1.47 percentage points on average, positive in 3 of 4 replicates. Worst-cell accuracy — the quantity we said in advance we would judge it on — moves by −0.23 points and wins in 2 of 4. The two summaries point in opposite directions, which is exactly the failure mode the pre-registration was written to expose. At this budget, on this task, forced adaptive acquisition does not reliably beat blind space-filling where it matters most. We report it because a negative result that was specified in advance is worth more than a positive one that was not.

8. What the agents actually built

A score on a frozen test set cannot distinguish two runs that arrive at the same error by entirely different reasoning. One might choose its adaptive points by ensemble disagreement and stop because validation error crossed the task’s threshold; another might relabel uniform sampling as “adaptive” and stop because it ran out of rounds. Those are different research behaviours and the same number.

Since that difference is much of what we wanted to measure, we measure it directly.

8.1. How it is measured

Two extractors run over each finished workspace. Both are **LLM-free and execution-free**: they parse the agent’s Python with the standard `ast` module and match patterns scoped to the role a file plays. The whole archive therefore re-extracts identically, for nothing, as often as we like.

Table 3. What each cell chose, and whether its code agrees with its prose. *Agr.* is the share of a cell’s replicates on its modal chain of methods — method reproducibility, which is not the same thing as score reproducibility. *Switch.* counts replicates whose acquisition criterion changed between rounds. *M-inf.* counts replicates whose point-choosing function actually calls the surrogate. The last column is the modal verdict of comparing the criterion each run names against the one its code computes.

Cell	Model	Acquisition	Agr.	Switch.	M-inf.	Code vs prose
L2-cursor	kernel ridge	<code>uncertainty_gp</code>	0.20	0/2	—	unverifiable 5/5
L2-dspy	kernel ridge	<code>residual_ucb</code>	0.10	1/9	10/10	partial 8, agree 2
L2-pi	kernel ridge	<code>residual_ucb</code>	0.20	1/5	—	unverifiable 5/5
L2-ursa	ridge / k.r.	<code>uncertainty</code>	0.38	0/3	2/5	partial 3, mismatch 2
L3-cursor	MLP (Torch)	<code>uncertainty_ensemble</code>	0.20	0/5	5/5	agree 4 , partial 1
L3-dspy	MLP (Torch)	<code>uncertainty_ensemble</code>	0.10	0/8	8/10	partial 9, mismatch 1
L3-pi	MLP (Torch)	<code>uncertainty_ensemble</code>	0.20	0/3	5/5	agree 3 , partial 2
L3-ursa	MLP (Torch)	<code>uncertainty</code>	0.17	0/2	4/5	partial 4, mismatch 1

Scope is load-bearing. A whole-workspace search for `subprocess` is true of nearly every run and tells you nothing, so each question is answered only from the code that plays the relevant role — the masking rule from `predict.py` and the solver wrapper, the self-test from whatever actually invokes `predict.py` as a subprocess — and files that are clearly not production code (tests, smoke checks, repro scripts) are excluded. Every answer carries `file:line` evidence, so any cell of the resulting table can be audited back to the line that produced it.

The vocabulary is anchored on the typed action space of the scripted levels, so that a scripted policy, a typed picker and a from-scratch agent all land in the same table. Anything unmatched stays **unknown** rather than being forced into a category. These are measurements placed beside the contract, never gates.

8.2. Method choice tracks the access level, not the substrate

Table 3 shows a split that no outcome metric reveals.

At L2, the primary model is **kernel ridge in 19 of the 25 replicates** with a classifiable primary — modal in three of the four cells, and tied first in the fourth. Kernel ridge is also what the library’s own model zoo selects in all six runs of the L0 arm. At L3 the primary model is **a hand-written PyTorch MLP in all 25**, without a single exception.

Acquisition splits the same way. L2 agents gravitate to residual-UCB and GP variance — the two strategies the L2 prompt names as available in the library — while L3 agents converge on ensemble disagreement, which no prompt mentions anywhere.

So: four independent agent frameworks, given a library, settle on its defaults or tie with them. Given nothing, all four rediscover a different method family, and two of them beat the library with it (Section 6.4). That is the mechanism behind the level effect of Section 6.2, and it is why we read that effect as anchoring rather than as the

library being poor code.

8.3. Method reproducibility is poor even where score reproducibility is good

Chain agreement — the share of a cell’s replicates that chose the same chain of methods — never exceeds 0.38, and is 0.10 in both ten-replicate cells. The same agent, given the same prompt and the same model, produces a different pipeline nearly every time.

L3-cursor is the sharpest case: five distinct method chains in five replicates, while its error stayed inside [0.062, 0.212] and every artifact was physically valid. A benchmark that replicates a cell and reports only the spread of scores would conclude the agent is consistent. It is consistently *good*. It is not doing the same thing twice, and if what you care about is whether an agent has a reliable method rather than a reliable outcome, those are very different findings.

8.4. Adaptivity is nominal: one fixed rule, executed n times

The task mandates up to three adaptive rounds whose points are chosen from the current model or data. In **2 of the 37** runs where the criterion could be classified across rounds did that criterion actually change between rounds. The rest ran one rule three times.

The crude failure — a stated criterion that is simply randomness — occurs in no cell at all. It is the subtler one, non-reactive adaptivity, that is near-universal. That is consistent with the control arm of Section 7.7 finding little for a reactive criterion to earn at this budget: the agents may be leaving nothing much on the table.

8.5. The code often does not match the prose, and a library hides it

Comparing the criterion each run *states* against the one its code *computes*, at the level of the criterion family, gives: **partial** 27, **unverifiable_from_code** 14, **agree** 9, **mismatch** 4, of the 54 runs with a verdict. **agree** is the modal verdict in only two cells, both L3. The four outright mismatches — stated and implemented criteria with nothing in common — fall in the two substrates that also produce every invalid artifact.

Seven L2 runs name a method in their README or report for which the code carries no evidence at all, against **zero such claims in any L3 cell**.

More awkward for anyone trying to evaluate an agent: L2-cursor and L2-pi are **unverifiable_from_code** in all ten of their runs, and whether their point selection is model-informed is simply unknown. Selection happens inside the library, so no function in the workspace chooses points at all. This is a structural consequence of giving an agent scaffolding rather than evasion by the agent — but it means that **the more scaffolding an evaluation provides, the less of the agent’s reasoning remains inspectable**. Where a chooser can be classified at all it usually does call the surrogate: 22 of 25 such runs at L3 and 12 of 15 at L2.

Relatedly, *ursa*’s rounds largely did not record validation error against baseline at the moment of choosing (median evidence-grounding 0.33 at L2 and 0.17 at L3, against

1.0 in the six other cells), so its stopping decisions cannot be falsified from its own artifacts.

8.6. Where the rest of the solution agrees, and where it only nearly does

Twelve further dimensions ask how each run solved the remainder of the task — who owns the solver’s single global environment, what insulates one solve from the next, how the triangular mesh reaches the fixed output grid, and so on. The full cross-comparison is in Appendix Appendix E.

One prohibition held outright: **no L3 run hand-built a mesh**, 26 of 26, which is the explicit anti-Delaunay constraint in the task text doing its job.

Three dimensions look unanimous at the level of each cell’s modal value, and a modal value hides the minority. The per-run counts, over the 51 runs whose workspace can be read, are: 50 enforce the pilot gate, 46 validate stored data before counting a solve — *three claim to and provide no evidence* — and 42 of the 43 with a readable answer keep the test set out of every model-fitting function. *One does not*.

The mandatory constraint most widely sidestepped is the deliverable self-test. The task requires re-running the documented `predict.py` command in a fresh process before declaring done. Twenty-four runs satisfy that by *importing* the module instead, against 22 that actually run it as a subprocess. Code that has not been re-run the way it is documented to be run is the precondition for most of the interface failures this benchmark has ever recorded — and it is measurable only because the question is asked of the code that plays that role, rather than of the workspace as a whole.

9. Limitations

- (i) **The primary statistical test was changed after we saw the data.** We pre-specified a sign test over four harness pairs, discovered it cannot return $p < 0.125$ at any effect size, and moved to a stratified rank test. All three tests agree in sign, so the *direction* of the result does not depend on the choice; the p -value does. All three are reported wherever the contrast appears (Section 6.2).
- (ii) **L3 receives two scaffolding instructions that L2 does not** — a mandatory meshing milestone and a warning against hand-built meshes. They are arguably necessary, since L3 agents must write solver plumbing that L2 agents delegate to a library call. But it means the L2/L3 contrast bundles library access together with a prompt difference, and should not be read as a pure library-access effect.
- (iii) **Replicate counts are unequal and partly determined by failure.** Two substrates have $n = 10$ and two $n = 5$. Denominators count attempts, so L2-dspy is 10/15 where among executed runs it is 10/10. Both conventions are defensible and the choice moves the pooled L2 pass rate materially, so we state ours wherever a pass rate appears.
- (iv) **One model, and the judge shares it.** Model control buys internal validity at the

cost of external: we cannot say these results hold for other models, and the harness effect may itself be model-dependent. Worse, the blind code review of Section 7 is performed by the same model that produced every run it grades. Self-preference cannot be excluded and no correction is applied, which is why those scores only ever corroborate the deterministic analysis and never stand alone.

- (v) **Wall-clock time is not a clean metric here.** Cells ran concurrently on one CPU-bound machine, so reported wall time reflects contention, and no separate executed-time measure was recorded, so a cell suspended by the operating system is indistinguishable from a slow one. Cost and turn counts are the efficiency axes we report.
- (vi) **Two accounting defects, disclosed.** One L3-ursa run was killed at a 1200s cap that later runs in the same cell did not face — the only such case. And two substrates reuse a run identifier across sibling replicate directories, so a join on that identifier reuses one replicate’s cost for its siblings; we report the directory-keyed total (\$197.78) and note the discrepancy (\$199.76).
- (vii) **Only two of the four ladder rungs are filled at this prompt version.** The L0 reference is three seeds, and no L1 typed-decision arm was re-run. The ladder is a contribution of the design; this campaign exercises half of it.
- (viii) **The physics is real but partial.** Fixed-boundary, a single profile family, five parameters. An agent that does well here has not been shown to do well on free-boundary equilibria, on time-dependent evolution, or on a device whose geometry is not a smooth D.

10. Conclusion

We built a benchmark in which an agent has to buy its own data, and then ran 59 agents through it on a single model.

Giving an agent a proven library made it worse. From-scratch agents beat library-assisted ones in all four substrates, by a pooled median rel- L_2 of 0.137 against 0.229, and the deterministic analysis of what they wrote says why: given a library, all four settle on its default model and its named acquisition strategies; given nothing, all four converge on a different family and two of them beat the library’s own score. Holding the model fixed, which framework wrapped it mattered more than how much help it was given. And the pipeline that uses no language model at all, at a matched budget and no cost, still beat five of the eight agent cells.

The finding we would most like to see travel is the one about verification. Thirty-one per cent of scored runs passed every machine-checkable gate while being physically invalid, and the single clearest example scored better than most of the corpus on the error metric while inventing a magnetic field in a vacuum. The runs that misreported their own performance were largely not the same runs, so neither measurement detects the other.

A benchmark that checks only whether the required files appeared will certify

work that is physically wrong and self-reported dishonestly, and it will do so silently. Separating *did it deliver*, *is it valid* and *did it tell the truth* costs three diagnostics computed from predictions the benchmark already holds. Asking the fourth question — *did it build what it said it built* — costs a parser. Between them they are the difference between measuring what an agent can do and measuring whether it filed the right paperwork.

Acknowledgments

The campaign ran on a single workstation with no external compute grant. We thank the Open FUSION Toolkit developers, whose solver and worked examples are the substrate this benchmark is built on, and the maintainers of the four agent frameworks evaluated here, none of whom were involved in or consulted about this evaluation.

Data availability statement

All data and code needed to reproduce this work are openly available. The benchmark — task specifications, the deliverable contract, the frozen evaluation grid and parameter set, every harness adapter and every analysis script — is at <https://github.com/hindyros/autotokamak>. The campaign artifacts analysed here (per-run records, agent transcripts, the aggregated result tables and the statistical output) are released with it. The ground-truth flux fields are regenerated deterministically from the committed parameter file and grid specification by a single command; they are not redistributed as a binary because every archived score is defined relative to a provenance-stamped rebuild.

Appendix A. Reproducing this work

```

pip install -e ".[ml,dev,harnesses]"
python -m autotokamak.bench freeze-testset      # rebuild the frozen ground truth

tools/campaign_guard.py preflight --harnesses "ursa dspy pi cursor" \
    --task benchmarks/tasks/L3_mini_v3.yaml
tools/run_campaign.sh --tag <tag> --reps 5 --parallel 6 --budget-usd 250 \
    --harnesses "ursa dspy pi cursor"

python tools/cost_report.py          --tag <tag>
python tools/aggregate_matrix.py     --tag <tag>    # main table + all three tests
python tools/render_paper_figures.py --tag <tag> --out <paper>/figures
python tools/render_physics_figures.py --tag <tag> --out <paper>/figures
python tools/render_paper_tables.py  --tag <tag> --out <paper>/tables

python tools/run_library_baseline.py --n-samples 450 --seeds 0 1 2    # the L0 arm
python tools/run_control_arm.py      --reps 4                        # Sec. 7.8
python tools/judge_code.py score     --tag <tag> --judge-provider openai \
    --samples 3 --per-cell 3

```

Every run records the task’s SHA-256, its prompt version, its model string and its replicate index. The frozen test set carries provenance attributes — the solver version, hashes of both the parameter file and the grid specification, and a scoring epoch — and is never regenerated in place; results are pooled only within one epoch. The statistics are seeded (bootstrap and both permutation tests use seed 20260917, with 10 000 and 20 000 resamples respectively), so the analysis reproduces exactly from the archived runs. Every figure and appendix table in this paper is generated from those artifacts by the scripts above rather than transcribed by hand.

The one thing that does *not* reproduce exactly is the agent runs themselves. They are stochastic, and the substrates are versioned third-party software that moves underneath us. That is why the campaign is replicated and reported with intervals rather than as point estimates.

Appendix B. Frozen assets

- `benchmarks/assets/eval_grid.json` — the evaluation grid: 64 points in $R \in [0.15, 0.80]$ m and 96 in $Z \in [-0.40, 0.40]$ m, both linearly spaced. Flux arrays are stored $(N, 96, 64)$, that is (N, n_Z, n_R) .
- `benchmarks/assets/test_params.json` — the 60 held-out parameter vectors, generated with seed 20260809.
- `benchmarks/assets/test_set.h5` — the ground-truth flux fields. All 60 solves converged; 80.46% of grid points are NaN, outside the plasma. Provenance attributes record the stamp time, the solver version (26.6), the SHA-256 of both the parameter file and the grid, the sweep configuration and its hash, and the scoring epoch (2026-08-18).
- `benchmarks/assets/prices.json` — the dated price table behind every derived cost figure, so that a published dollar amount is traceable to the rates it was computed with.
- `benchmarks/assets/scaling_report.json` — the task-characterisation sweep behind Figure 4 and Appendix Appendix C.
- `benchmarks/tasks/*.yaml` — the task specifications.

Appendix C. Task characterisation in full

From a direct sweep independent of any agent: a pool of 9993 converged solves with a 1499-case held-out test split. Error is reported as a ratio to the mean-predictor baseline, so 1.0 means “no better than predicting the average field”. Summarised as Figure 4.

Dimensionality reduction bottoms out near $k \approx 24$ PCA components (ratio 0.293) and degrades slightly by $k = 32$, confirming the field is low-rank but not as low-rank as normalised flux would be. Retaining physical webers costs rank and is still the right choice, because per-sample normalisation destroys the I_p dependence entirely.

Table C1. Scaling of surrogate error with training-set size. “Best ratio” is relative L_2 divided by the mean-predictor baseline, minimised over the model zoo. The winning family changes at $N = 400$.

N	50	100	200	400	800	1600	3200	6400
best ratio	0.575	0.526	0.485	0.465	0.369	0.303	0.278	0.268
R^2	0.804	0.840	0.865	0.876	0.922	0.948	0.956	0.959

Table C2. Error ratio by sampling design at matched budget.

N	space-filling	random	clustered
100	0.550	0.699	0.963
400	0.358	0.438	0.873
1600	0.289	0.309	0.808

Table C3. Out-of-distribution generalisation. The transfer directions, in bold, collapse to the mean predictor or worse.

train $\rightarrow$ test	ratio	R^2
low $\rightarrow$ low	0.376	0.927
high $\rightarrow$ high	0.401	0.919
low $\rightarrow$ high	0.943	0.280
high $\rightarrow$ low	0.937	-0.097

Appendix D. The task text, and how the two levels differ

The L3 problem statement is reproduced verbatim below. The L2 text is byte-identical except for the ten passages that follow, every one of which concerns access level: the title, the description of the solver, the three pipeline steps where the library would be used, the workspace and environment bullets, the advice about the solver’s single global environment, and the constraints block. The parameter ranges, the budget, the stopping criterion, the evaluation interface, all five deliverables, the instruction not to fabricate metrics, and the three process gates are identical.

Appendix D.1. The L3 problem statement

TITLE

Build, end to end, a machine-learning surrogate for fixed-boundary Grad-Shafranov tokamak equilibria - including a data-generation campaign with adaptive (active) sampling - from scratch.

THE TASK

TokaMaker, part of the pip-installed ‘OpenFUSIONToolkit’ Python package, solves the Grad-Shafranov equation for the poloidal flux $\psi(R, Z)$ of a tokamak plasma. You will use it as your ground-truth physics solver.

Consider fixed-boundary equilibria of D-shaped plasmas described by 5 input parameters, with TokaMaker’s default pressure/current profile shapes:

```

r0   : major radius      [m]   in [0.35, 0.55]
a    : minor radius      [m]   in [0.10, 0.20]
kappa : elongation       [-]   in [1.0, 1.6]
delta : triangularity     [-]   in [0.0, 0.4]
Ip    : plasma current    [A]   in [80e3, 200e3]
(plasma centered at z0 = 0)

```

Your job, in full:

1. Learn how to drive TokaMaker for fixed-boundary solves (mesh a D-shaped last-closed-flux-surface, solve, extract psi). Choose your own numerical resolution; a single solve should take seconds, not minutes.
2. Design and run a DATA-GENERATION CAMPAIGN structured as an initial space-filling design of AT MOST 150 solves followed by UP TO 3 ADAPTIVE (ACTIVE-SAMPLING) ROUNDS of EXACTLY 100 new solver runs each (failed/unconverged solves count against a round's 100). A one-shot space-filling or random design alone is NOT sufficient: each round's 100 new solver evaluations must be chosen ADAPTIVELY, informed by the current model or data (your strategy - e.g. uncertainty, model error, disagreement, anything you can justify). STOPPING CRITERION: keep a held-out VALIDATION set (separate from the final test set of step 4, which must never influence the campaign); after each round, retrain your surrogate and stop the campaign early once its validation error is reduced by at least 70% relative to the mean-predictor baseline (a model that always predicts the training-set mean psi map), i.e. $\text{val_error} \leq 0.30 * \text{baseline_error}$. Otherwise continue to the 3-round cap. Log every acquisition decision (which point, when, why) and each round's validation error vs baseline to an acquisition log file.
3. Train an ML surrogate mapping the 5 parameters to the PHYSICAL poloidal flux $\psi(R, Z)$ in Webers (not a normalized flux). Model family, input/output representation, and training procedure are entirely your choice.
4. Evaluate honestly on a held-out test set of AT LEAST 20 solver runs at uniformly random parameters (fixed recorded seed), generated independently and never used for training, model selection, or acquisition.
5. Report results and analyze whether your adaptive sampling actually helped (e.g. versus the model trained on the initial design only, or any comparison you can defend).

EVALUATION INTERFACE (fixed, so results are comparable across methods)

Predictions and metrics are defined on this fixed rectangular grid:

R : 64 points, uniform in [0.15, 0.80] m

Z : 96 points, uniform in [-0.40, 0.40] m

psi arrays are shaped (96, 64) = (nZ, nR). Grid points outside the plasma boundary carry NaN; all metrics are computed only where the ground truth is finite.

Required metric: per-sample relative L2 error

$\text{err}_i = ||\psi_{\text{pred}_i} - \psi_{\text{true}_i}||_2 / ||\psi_{\text{true}_i}||_2$

over the finite grid points. Report mean, median, and 90th percentile over the test set, and the same numbers for a trivial baseline (predict the training-set mean psi) so the surrogate's value is evident.

DELIVERABLES (all in your current working directory; internal file formats and code organization are your choice)

1. Your pipeline code, runnable end to end by one documented entry command.
2. The dataset(s) you generated, plus the acquisition log.
3. The trained model artifact, plus a prediction script with EXACTLY this interface:

```
python predict.py --input params.json --output pred.npz
```

 where params.json is a JSON list of objects with keys
 r0, a, kappa, delta, Ip, and pred.npz contains:

```
psi : (N, 96, 64) float64      predicted flux [Wb], NaN outside plasma
R    : (64,) float64           grid R coordinates [m]
```

- Z : (96,) float64 grid Z coordinates [m]
4. report.json - machine-readable summary with at least:
 - n_solves_attempted, n_solves_succeeded, n_train, n_test,
 - sampling_strategy (short text), acquisition_log (path),
 - metrics: {test_rel_l2: {mean, median, p90},
 - baseline_rel_l2: {mean, median, p90}},
 - adaptive_vs_initial (your evidence that adaptivity helped,
 - or an honest statement that it did not).
 5. README.md - your design, why you chose it, how to run, results.

You MUST actually run the full campaign and training within this session and verify every deliverable exists before declaring done. Report failures honestly in report.json/README; do not fabricate metrics.

ENVIRONMENT FACTS

- * Your current working directory IS your workspace. Create all files here. Do not read, copy, or import code from the surrounding repository. In particular the Python environment contains an installed package called 'autotokamak' that solves a related problem - importing it, reading its source, or reusing its artifacts is FORBIDDEN: this is a from-scratch capability test, and only 'OpenFUSIONToolkit' may be used for physics.
- * 'python' resolves to a Python 3.11 venv that already has: numpy, scipy, h5py, matplotlib, pyyaml, scikit-learn, torch, optuna, OpenFUSIONToolkit. Nothing may be installed.
- * A read-only './OpenFUSIONToolkit/' symlink to the full OFT source tree is present in your workspace. BEFORE writing any solver or meshing code, study the worked TokaMaker examples under 'OpenFUSIONToolkit/src/examples/TokaMaker/' (the 'fixed_boundary/' example is directly relevant) and follow the workflow they demonstrate rather than inventing your own solver-facing plumbing.
- * TokaMaker limitation: only ONE 'OFT_env' may ever be created per Python process. Structure your code so the environment object is created once and reused across all solves in a process, or run solves in child processes.

CONSTRAINTS

- * MESHING MILESTONE (mandatory, before any campaign code): locate the fixed-boundary example under 'OpenFUSIONToolkit/src/examples/TokaMaker/', reproduce its mesh-and-solve workflow as a standalone script in your workspace, run it, and verify it produces finite psi values. Only then generalize it to your parameterized D-shape. Do not design your own meshing approach when the examples demonstrate one.
- * Do NOT construct the computational mesh yourself (e.g. via scipy Delaunay) and pass raw triangles to the solver - TokaMaker validates mesh topology and will reject hand-built triangulations. Use the meshing route the OFT examples use.
- * PILOT GATE (mandatory): before the campaign, run a pilot of at least 20 solves at random parameters. If the success rate is below 50%, STOP, diagnose the per-failure error messages, fix, and re-run the pilot. Never start the 100-solve rounds while the pilot is failing.
- * STORAGE VALIDATION GATE (mandatory): a solver run may be counted as successful ONLY after its stored artifact has been re-loaded from disk and verified: the saved psi grid must contain a nonzero fraction of finite values. Apply this check to every batch you write. A stored all-NaN psi grid is a FAILED solve regardless of solver exit status, and your reported success counts must be based on validated stored artifacts, not solver return codes.
- * DELIVERABLE SELF-TEST (mandatory, before declaring done): in a fresh Python process, run the exact documented command for each deliverable - including 'python predict.py --input <sample params> --output <out>' - and verify the output loads with the required shapes and finite values. If any deliverable fails its own interface test, fix it and re-test. Code you have not re-run after your last edit does not count as verified.
- * Do not run 'pip install' or modify the environment.

- * Do not call ‘git’.
- * Do not prompt for user input; the task is fully autonomous.
- * If errors arise, debug, fix your code, and re-run. Stay within the campaign structure above (at most 3 adaptive rounds x 100 solves, plus the initial design and validation/test sets) across ALL attempts that reach the solver.

Appendix D.2. Every difference from the L2 statement

```

--- L2_mini_v3
+++ L3_mini_v3
@@ -4,2 +4 @@
- campaign with adaptive (active) sampling - using the installed
- ‘autotokamak’ platform library wherever it helps.
+ campaign with adaptive (active) sampling - from scratch.
@@ -10,2 +9,2 @@
- psi(R, Z) of a tokamak plasma. It is your ground-truth physics
- solver, and the installed ‘autotokamak’ package already wraps it.
+ psi(R, Z) of a tokamak plasma. You will use it as your ground-truth
+ physics solver.
@@ -25,4 +24,3 @@
- 1. Drive the solver for fixed-boundary solves. You may use
- ‘autotokamak.data.sweep.run_sweep’ (config-driven sweeps; also
- accepts an explicit (N, 5) sample matrix), or go through
- ‘autotokamak.core’ directly. A single solve should take
+ 1. Learn how to drive TokaMaker for fixed-boundary solves (mesh a
+ D-shaped last-closed-flux-surface, solve, extract psi). Choose
+ your own numerical resolution; a single solve should take
@@ -36,13 +34,13 @@
- chosen ADAPTIVELY, informed by the current model or data
- (‘autotokamak.data.acquire’ implements residual-UCB and PCA-GP
- variance acquisition you may reuse, or bring your own).
- STOPPING CRITERION: keep a held-out VALIDATION set (separate
- from the final test set of step 4, which must never influence
- the campaign); after each round, retrain your surrogate and
- stop the campaign early once its validation error is reduced
- by at least 70% relative to the mean-predictor baseline (a
- model that always predicts the training-set mean psi map),
- i.e. val_error <= 0.30 * baseline_error. Otherwise continue to
- the 3-round cap. Log every acquisition decision (which point,
- when, why) and each round’s validation error vs baseline to an
- acquisition log file.
+ chosen ADAPTIVELY, informed by the current model or data (your
+ strategy - e.g. uncertainty, model error, disagreement,
+ anything you can justify). STOPPING CRITERION: keep a held-out
+ VALIDATION set (separate from the final test set of step 4,
+ which must never influence the campaign); after each round,
+ retrain your surrogate and stop the campaign early once its
+ validation error is reduced by at least 70% relative to the
+ mean-predictor baseline (a model that always predicts the
+ training-set mean psi map), i.e. val_error <= 0.30 *
+ baseline_error. Otherwise continue to the 3-round cap. Log
+ every acquisition decision (which point, when, why) and each
+ round’s validation error vs baseline to an acquisition log
+ file.
@@ -51,4 +49,2 @@
- ‘autotokamak.surrogate’ provides dataset loading, PCA
- reduction, a model zoo (GP, kernel ridge, polynomial ridge,
- MLP), and Optuna search; model family and representation
- remain entirely your choice.
+ Model family, input/output representation, and training
+ procedure are entirely your choice.
@@ -107 +103,6 @@
- here.
+ here. Do not read, copy, or import code from the surrounding
+ repository. In particular the Python environment contains an
+ installed package called ‘autotokamak’ that solves a related
+ problem - importing it, reading its source, or reusing its
+ artifacts is FORBIDDEN: this is a from-scratch capability test,

```

```

+ and only 'OpenFUSIONToolkit' may be used for physics.
@@ -110,8 +111,8 @@
- optuna, OpenFUSIONToolkit, and the 'autotokamak' platform library
- (importable; using it is ALLOWED and RECOMMENDED - that is the
- point of this condition). Nothing may be installed.
- * Useful autotokamak entry points: 'autotokamak.data.sweep.run_sweep'
- (sweep or explicit-matrix solves → HDF5),
- 'autotokamak.data.acquire' (acquisition strategies),
- 'autotokamak.surrogate.dataset/reduce/zoo/optuna_search/automl_loop'
- (loading, PCA, model factories, search). Read their docstrings.
+ optuna, OpenFUSIONToolkit. Nothing may be installed.
+ * A read-only './OpenFUSIONToolkit/' symlink to the full OFT
+ source tree is present in your workspace. BEFORE writing any
+ solver or meshing code, study the worked TokaMaker examples
+ under 'OpenFUSIONToolkit/src/examples/TokaMaker/' (the
+ 'fixed_boundary/' example is directly relevant) and follow the
+ workflow they demonstrate rather than inventing your own
+ solver-facing plumbing.
@@ -119,4 +120,3 @@
- Python process. 'autotokamak.core.solver.make_solver' handles
- reuse via its 'env=' argument; if you bypass autotokamak, create
- the environment object once and reuse it, or run solves in child
- processes.
+ Python process. Structure your code so the environment object is
+ created once and reused across all solves in a process, or run
+ solves in child processes.
@@ -124,0 +125,11 @@
+ * MESHING MILESTONE (mandatory, before any campaign code): locate
+ the fixed-boundary example under
+ 'OpenFUSIONToolkit/src/examples/TokaMaker/', reproduce its
+ mesh-and-solve workflow as a standalone script in your workspace,
+ run it, and verify it produces finite psi values. Only then
+ generalize it to your parameterized D-shape. Do not design your
+ own meshing approach when the examples demonstrate one.
+ * Do NOT construct the computational mesh yourself (e.g. via scipy
+ Delaunay) and pass raw triangles to the solver - TokaMaker
+ validates mesh topology and will reject hand-built
+ triangulations. Use the meshing route the OFT examples use.
@@ -147,2 +157,0 @@
- * Do not modify the installed autotokamak/OpenFUSIONToolkit
- packages; use them as libraries.

```

Appendix E. Solution-shape cross comparison

Twelve questions, each answered from the code that plays the relevant role, with file:line evidence recorded per run. Cells are abbreviated by level and harness: 2c is L2-cursor, 3u is L3-ursa, and so on. A dash means the dimension is not answerable for that cell, which at L2 generally means the library owns that part of the pipeline and no workspace code implements it.

Dimension	Cells	Modal value
code_shape	2c, 2p	single_script
	2d	14_modules
	2u	15_modules
	3c	8_modules
	3d	11_modules
	3p	9_modules
	3u	7_modules
entry_point	2c, 2p, 3p	2_runnable_scripts
	2d, 3u	10_runnable_scripts
	2u	9_runnable_scripts
	3c	3_runnable_scripts
	3d	7_runnable_scripts
grid_mapping	2c, 2d, 2p, 2u	—
	3c, 3d, 3p, 3u	solver_field_eval
mask_rule	2c, 2p, 3c, 3d, 3p	lcfs_polygon
	2d, 2u	training_valid_mask
	3u	lcfs_polygon+training_valid_mask
mesh_route	2c, 2u, 3c, 3d, 3p, 3u	oft_gs_domain
	2d, 2p	—
oft_env_strategy	2c, 2d, 2p, 2u	—
	3c, 3d, 3p, 3u	singleton_reused
self_test	2c, 2d, 2u, 3p, 3u	subprocess_reruns_predict
	2p, 3c, 3d	imported_not_subprocessed
solve_isolation	2c, 2d, 2p, 2u, 3c, 3d, 3u	in_process_serial
	3p	subprocess_per_batch
storage	2c, 2p, 3c, 3p	npz_per_solve+pickle
	3d, 3u	npz_per_solve+pickle+index
	2d	npz_per_solve+hdf5
	2u	npz_per_solve+hdf5+index
leakage_guard	2c, 2d, 2p, 2u, 3c, 3d, 3p, 3u	test_absent_from_fitting_functions
pilot_gate	2c, 2d, 2p, 2u, 3c, 3d, 3p, 3u	enforced
storage_validation	2c, 2d, 2p, 2u, 3c, 3d, 3p, 3u	reload_and_check_finite

Appendix F. Glossary of canonical values

Every code printed in Table 3 and Appendix Appendix E. A code the reader has to guess at is not a measurement.

N_modules

The pipeline is spread over N source files, none of them dominant; contrast `single_script`.

N_runnable_scripts

N files are runnable as programs. More than one means the documented single command is really a sequence.

adaptive_in_name_only

Every stated criterion names nothing but randomness: the round was labelled adaptive and was not.

agree

The criterion stated in the log/report is the one the code computes.

chain_agreement

Share of a cell's replicates that chose the modal chain of methods — method reproducibility, which is not the same as score reproducibility.

criterion_switched

The acquisition criterion CHANGED between rounds — logic reacting to what it measured, rather than one fixed rule executed n times.

deep_ensemble

Several independently initialised models, averaged.

enforced

A pilot runs and its success rate is compared against a threshold before the campaign proceeds — the task's 50% pilot gate.

evidence_grounded

The round's validation error against baseline was recorded, so its choice can be checked against what was known at the time. An ungrounded round's reasoning is unfalsifiable.

feasibility

Candidate choice weighted by whether solves there are expected to converge, steering away from regions with failed solves.

gradient_sensitivity

Points where the response is steep or curving — a sensitivity or Jacobian-driven criterion.

hdf5

One HDF5 file holding many samples. Scales best and carries attributes/provenance, at the price of concurrent-write care.

imported_not_subprocessed

predict.py is referenced but only imported or described, never run as the documented command. Code that works in the session's process and fails in a clean one is exactly what this misses.

in_process_serial

Solves run one after another inside the driving process. Simplest, and the slowest: no parallelism, and one bad solve can poison the shared solver state.

index

An index/manifest (CSV or JSONL) sits beside the arrays, listing every sample and its status. What makes a campaign resumable and auditable rather than a directory to be re-scanned.

kernel_ridge

Kernel ridge regression — a GP's mean without its variance.

lcfs_polygon

A grid point is inside the plasma iff it lies within the LCFS polygon — point-in-polygon against the D-shape boundary, whatever the helper is called (contains_points, _points_in_poly, inside_lcfs_mask). Geometric, per sample, and independent of the surrogate's own output.

lhs

Latin hypercube: stratified in every dimension at once.

mc_dropout

Dropout left active at inference and sampled repeatedly — an ensemble’s spread at one model’s training cost.

mismatch

The stated criterion and the implemented one have nothing in common. The strongest available signal that a run’s account of itself is wrong.

mlp_torch

A hand-written PyTorch MLP. The most common choice in this corpus, usually mapping 5 parameters to PCA coefficients.

model_informed

The function that chooses points actually calls the surrogate. A model-derived criterion that never does is not one, whatever it is named.

npz_per_solve

One compressed .npz per solved sample.

oft_gs_domain

Meshed through OFT’s own API (`gs_Domain / define_region / build_mesh`) as the worked examples do. What the task’s meshing milestone asks for, and what TokaMaker’s topology checks accept.

partial

Stated and implemented criteria overlap but do not match: one side names a component the other never does — see `claimed-but-not-computed`.

pickle

Python-pickled objects (`pickle/joblib/torch.save`) — usually the model artifact rather than the dataset.

prose_only

A method named in the README or `report.json` that the code never evidences.

random

Points drawn uniformly at random. As the CANDIDATE POOL this is normal and harmless; as the selection criterion it means the round was adaptive in name only.

reload_and_check_finite

A function both re-loads a written artifact and checks its finite fraction — the task’s storage-validation gate, actually implemented. This is the check that catches an all-NaN grid stored as a success.

residual_uch

Points are targeted where the model is MEASURABLY wrong — an error model fit on residuals or out-of-fold error, often with a UCB trade-off between exploiting known weakness and exploring. Uses measurement rather than the model’s self-assessment, which can be confidently wrong.

round_cap

Stop because the 3-round cap was reached.

single_script

The whole pipeline is one or two files.

singleton_reused

One `OFT_env` is built once and handed out on every call — a cached module global, a class guarding construction, or an `lru_cache`. Honours OFT’s one-env-per-process limit in the simplest way, at the cost of keeping every solve in one process: a solver crash takes the campaign with it.

sobol

A Sobol low-discrepancy sequence — quasi-random, extensible.

solver_field_eval

`psi` is read on the frozen 64x96 grid through the solver’s own finite-element interpolator (`get_field_eval`). Evaluates the FE basis directly, so no second interpolation error is introduced.

space_filling

Purely geometric coverage of the input box: farthest-point/maximin, distance to the existing training set, k-means. Needs no model at all — which makes it robust, and makes it NOT adaptive in the sense of reacting to what was learned.

stop_threshold_met

The campaign ended with the task’s 70% criterion satisfied — the intended way to finish early.

stop_unexplained

The campaign ended and no per-round validation error was recorded, so why it stopped cannot be read from the artifacts at all.

stop_without_threshold

The campaign ended although validation error had NOT reached 0.30x baseline: the round cap, a budget, or a timeout ended it, not the stopping rule.

subprocess_per_batch

The driver shells out (usually to its own script) to run a batch. Insulates solver state between batches and survives a hard crash; pays process start-up and serialises through files.

subprocess_reruns_predict

Something in the workspace invokes `predict.py` as a subprocess with the contract CLI — the task’s deliverable self-test, in a fresh process, which is the only way to catch an import-order or global-state dependency.

test_absent_from_fitting_functions

No function that calls `.fit()/backward()/optimizer.step()` references a test path or array. Necessary, not sufficient: the task also forbids the test set influencing model SELECTION and ACQUISITION, which this check cannot see.

top_k

The highest-scoring candidates are taken (`argsort/topk`).

training_valid_mask

The predictor masks with a valid-pixel mask derived from the training data (e.g. pixels finite in training). Cheap and stable, but the mask cannot adapt to a geometry unlike those seen in training — the same family of defect as the run that passed 9/9

gates while predicting plasma in vacuum.

uncertainty

Points are ranked by the model’s own predictive spread, without the record saying where that spread comes from. The family term; the two entries below are its specific forms.

uncertainty_ensemble

Spread across an ENSEMBLE of models (deep ensemble, bootstrap, or MC-dropout samples) — ‘query by committee’. Needs no probabilistic model, and its quality depends entirely on the members disagreeing for real rather than sharing an initialisation.

uncertainty_gp

Posterior variance of a Gaussian process (or an acquisition built on it, e.g. expected improvement). Principled and calibrated where the GP’s kernel assumptions hold; costly as the dataset grows.

uniform_random

Independent uniform draws, with no stratification.

unverifiable_from_code

No function that chooses points could be classified, so the stated criterion can be neither confirmed nor contradicted.

val_threshold_70pct

The task’s own rule: stop once validation error is at most 0.30x the mean-predictor baseline, i.e. a 70% reduction.

References

References

- [1] C. Hansen, I. G. Stewart, D. Burgess, M. Pharr, S. Guizzo, F. Logak, A. O. Nelson, and C. Paz-Soldan. TokaMaker: An open-source time-dependent Grad-Shafranov tool for the design and modeling of axisymmetric fusion devices. *Computer Physics Communications*, 298:109111, 2024. doi: 10.1016/j.cpc.2024.109111. arXiv:2311.07719.
- [2] H. Grad and H. Rubin. Hydromagnetic equilibria and force-free fields. In *Proceedings of the Second United Nations International Conference on the Peaceful Uses of Atomic Energy*, volume 31, pages 190–197. United Nations, Geneva, 1958.
- [3] V. D. Shafranov. On magnetohydrodynamical equilibrium configurations. *Soviet Physics JETP*, 6:545–554, 1958.
- [4] Plamen G. Krastev. Millisecond-scale neural operator surrogates for double-null free-boundary Grad-Shafranov equilibria. *arXiv preprint arXiv:2608.05555*, 2026. Harvard University; submitted 6 August 2026.
- [5] M. Grandin, A. Mele, C. Galperti, D. Gonzales Castineiras, C. Heiß, and A. Cenedese. Experimental validation of a fast control-oriented, physics-informed

- surrogate model for plasma equilibrium reconstruction in the TCV tokamak. *arXiv preprint arXiv:2606.09487*, 2026. with the TCV team and the EUROfusion Tokamak Exploitation team; submitted 8 June 2026.
- [6] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *International Conference on Learning Representations (ICLR)*, 2024.
 - [7] Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, Minlie Huang, Yuxiao Dong, and Jie Tang. AgentBench: Evaluating LLMs as agents. In *International Conference on Learning Representations (ICLR)*, 2024.
 - [8] Jun Shern Chan, Neil Chowdhury, Oliver Jaffe, James Aung, Dane Sherburn, Evan Mays, Giulio Starace, Kevin Liu, Leon Maksin, Tejal Patwardhan, Lilian Weng, and Aleksander Madry. MLE-bench: Evaluating machine learning agents on machine learning engineering. In *International Conference on Learning Representations (ICLR)*, 2025.
 - [9] Qian Huang, Jian Vora, Percy Liang, and Jure Leskovec. MAgentBench: Evaluating language agents on machine learning experimentation. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *PMLR*, pages 20271–20309, 2024.
 - [10] Liqiang Jing, Zhehui Huang, Xiaoyang Wang, Wenlin Yao, Wenhao Yu, Kaixin Ma, Hongming Zhang, Xinya Du, and Dong Yu. DS Bench: How far are data science agents from becoming data science experts? In *International Conference on Learning Representations (ICLR)*, 2025.
 - [11] Ziru Chen, Shijie Chen, Yuting Ning, Qianheng Zhang, Boshi Wang, Botao Yu, Yifei Li, Zeyi Liao, Chen Wei, Zitong Lu, Vishal Dey, Mingyi Xue, Frazier N. Baker, Benjamin Burns, Daniel Adu-Ampratwum, Xuhui Huang, Xia Ning, Song Gao, Yu Su, and Huan Sun. ScienceAgentBench: Toward rigorous assessment of language agents for data-driven scientific discovery. In *International Conference on Learning Representations (ICLR)*, pages 12138–12194, 2025.
 - [12] Bodhisattwa Prasad Majumder, Harshit Surana, Dhruv Agarwal, Bhavana Dalvi Mishra, Abhijeetsingh Meena, Aryan Prakhar, Tirth Vora, Tushar Khot, Ashish Sabharwal, and Peter Clark. DiscoveryBench: Towards data-driven discovery with large language models. In *International Conference on Learning Representations (ICLR)*, 2025. Allen Institute for AI.
 - [13] Ludovico Mitchener, Jon M. Laurent, Alex Andonian, Benjamin Tenmann, Siddharth Narayanan, Geemi P. Wellawatte, Andrew White, Lorenzo Sani, and Samuel G. Rodrigues. BixBench: A comprehensive benchmark for LLM-based agents in computational biology. *arXiv preprint arXiv:2503.00096*, 2025. FutureHouse.

- [14] Zachary S. Siegel, Sayash Kapoor, Nitya Nadgir, Benedikt Stroebel, and Arvind Narayanan. CORE-Bench: Fostering the credibility of published research through a computational reproducibility agent benchmark. *Transactions on Machine Learning Research (TMLR)*, 2025.
- [15] Giulio Starace, Oliver Jaffe, Dane Sherburn, James Aung, Jun Shern Chan, Leon Maksin, Rachel Dias, Evan Mays, Benjamin Kinsella, Wyatt Thompson, Johannes Heidecke, Amelia Glaese, and Tejal Patwardhan. PaperBench: Evaluating AI’s ability to replicate AI research. *arXiv preprint arXiv:2504.01848*, 2025.
- [16] Hjalmar Wijk, Tao Lin, Joel Becker, Sami Jawhar, Neev Parikh, Thomas Broadley, Lawrence Chan, Michael Chen, Josh Clymer, Jai Dhyani, Elena Elicheva, Katharyn Garcia, Brian Goodrich, Nikola Jurkovic, Holden Karnofsky, Megan Kinniment, Aron Lajko, Seraphina Nix, Lucas Sato, William Saunders, Maksym Taran, Ben West, and Elizabeth Barnes. RE-Bench: Evaluating frontier AI R&D capabilities of language model agents against human experts. *arXiv preprint arXiv:2411.15114*, 2024. METR.
- [17] Burr Settles. Active learning literature survey. Computer Sciences Technical Report 1648, University of Wisconsin–Madison, 2009.
- [18] Kathryn Chaloner and Isabella Verdinelli. Bayesian experimental design: A review. *Statistical Science*, 10(3):273–304, 1995. doi: 10.1214/ss/1177009939.
- [19] M. D. McKay, R. J. Beckman, and W. J. Conover. A comparison of three methods for selecting values of input variables in the analysis of output from a computer code. *Technometrics*, 21(2):239–245, 1979.
- [20] Jonathan Bragg, Mike D’Arcy, et al. AstaBench: Rigorous benchmarking of AI agents with a scientific research suite. *arXiv preprint arXiv:2510.21652*, 2025. Allen Institute for AI.
- [21] Yunbei Zhang, Janet Wang, Yingqiang Ge, Weijie Xu, Jihun Hamm, and Chandan K. Reddy. Stop comparing LLM agents without disclosing the harness. *arXiv preprint arXiv:2605.23950*, 2026.
- [22] Michael Grosskopf, Nathan Debardeleben, Russell Bent, Rahul Somasundaram, Isaac Michaud, Arthur Lui, Alexius Wadell, Warren D. Graham, Golo A. Wimmer, Sachin Shivakumar, Joan Vendrell Gallart, Harsha Nagarajan, and Earl Lawrence. URSA: The universal research and scientific agent, 2025. Los Alamos National Laboratory. Software: <https://github.com/lanl/ursa>.
- [23] Harrison Chase. LangChain. <https://github.com/langchain-ai/langchain>, 2022. Software; citation metadata from the repository’s CITATION.cff, dated 2022-10-17.
- [24] Omar Khattab, Arnav Singhvi, Paridhi Maheshwari, Zhiyuan Zhang, Keshav Santhanam, Sri Vardhamanan, Saiful Haq, Ashutosh Sharma, Thomas T. Joshi, Hanna Moazam, Heather Miller, Matei Zaharia, and Christopher Potts. DSPy: Compiling declarative language model calls into self-improving pipelines. In *International Conference on Learning Representations (ICLR)*, 2024.

- [25] Lakshya A. Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziemis, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J. Ryan, Meng Jiang, Christopher Potts, Koushik Sen, Alexandros G. Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and Omar Khattab. GEPA: Reflective prompt evolution can outperform reinforcement learning. In *International Conference on Learning Representations (ICLR)*, 2026. Oral. arXiv:2507.19457.
- [26] Edwin B. Wilson. Probable inference, the law of succession, and statistical inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927. doi: 10.1080/01621459.1927.10502953.
- [27] B. Efron. Bootstrap methods: Another look at the jackknife. *The Annals of Statistics*, 7(1):1–26, 1979. doi: 10.1214/aos/1176344552.
- [28] P. H. van Elteren. On the combination of independent two-sample tests of wilcoxon. *Bulletin of the International Statistical Institute*, 37:351–361, 1960.